University of Passau

Faculty of Computer Science and Mathematics

Chair for IT Security

UNIVERSITY OF PASSAU

Master Thesis in Informatics

Secure Bootstrapping and Post-Compromise Security in IoT

submitted by

Hossam Hamed

1. Examiner:	Prof. Dr. Joachim Posegga
2. Examiner:	Prof. Dr. Jorge Cuellar
Supervisor:	Fehlende Betreuerangabe
Date:	January 27, 2022

Contents

Abstract

Please write a short abstract summarizing your work.

Acknowledgments

I would first like to thank . . .

List of Figures

List of Tables

Acronyms

PKI Public Key Infrastructure

ABC Automated Border Control

DAS Document Authentication System

BVS Biometric Verification System

CSI Central Systems Interface

BGMS Border Guard Maintenance System

VMS Visa Management System

RTP Registered Traveler Program

EEMS Entry-Exit Management System

TCP Transmission Control Protocol

UDP User Datagram Protocol

IoT Internet of Things

TLS Transport Layer Security

DTLS Datagram Transport Layer Security

IETF Internet Engineering Task Force

CoAP Constrained Application Protocol

REST Representational state transfer

IDevID Initial Device Identifier

LDevID Locally Significant Device Identifier

P2P Peer-to-Peer

BRSKI Bootstrapping Remote Secure Key Infrastructure

SZTP Secure Zero Touch Provisioning

OOB Out-Of-Band

EAP-TLS Extensible Authentication Protocol TLS

GBA Generic Bootstrapping Architecture

EST Enrollment over Secure Transport

CA Certification Authority

MASA Manufacturer Authorized Signing Authority

RA Registration Authority

EE End Entity

DoS Denial of Service

SCRAM Salted Challenge Response Authentication Mechanism

TOFU Trust on first use

NEA Network Endpoint Assessment

CMP Certificate Managment Protocol

NMS Network Managment System

RESTCONF Representational State Transfer Configuration

OTR Off-the-Record

X3DH Extended Triple Diffie-Hellmann

KDF Key Derivation Function

DH Diffie-Hellmann

1 Introduction

1.1 Introduction and Motivation

Your thesis should be motivated in this chapter. Also outline the research gap to existing work.

1.2 Research Questions

Write down and explain your research questions

1.3 Structure of the Thesis

Explain the structure of your thesis.

2 Background

This chapters gives an overview over concepts important for the protocols and algorithms discussed in further chapters.

2.1 Certificate

2.2 Bootstrapping

2.3 Voucher Artifact

-nonced and nonceless

2.4 Certificate Enrollment

2.5 Protocol Formal Modeling

2.6 Encryption

2.6.1 Symmetric Encryption

2.6.2 Asymmetric Encryption

2.7 Key Derivation Function (KDF)

2.8 Threat Models

2.9 Secure Messaging

2.9.1 properties

2.9.2 OTR

3 Use Cases

3.1 ABC

International passenger traffic has been monitored to continuously increase over the years. Since air transport is one of the most convenient form of long distance transport among passengers, the above claim is indicated by the Airports Council International's Passenger Traffic Summary[1]. Despite the current shock due to the Covid-19 pandemic to the air transport industry, the aviation industry has shown resilience over the past decades and is predicted to recover and further grow by 4% over the course of the next 20 years[2] international border crossing points are forecasted to face more traffic than they already are. Implying the requirement for a solution to the Border control process to mitigate longer queues and waiting times, or the need to hire and train more personnel to conduct the task.

bio-metric passports are widely issued by more than 150 countries and regions, as of 2020[3]. Such passports contain a contactless smart chip that stores the passport holder's bio-metric information to authenticate his identity at passport control stations. The chips may contain one or more of the following bio-metric data: Iris, Fingerprint, and/or Facial features. In recent years, ABC systems, or E-Gates, started rolling out in airports leveraging the wide deployment of bio-metric passports. They intend to automate the process and improve the management and control of travel flows, serving passengers in a shorter time and reduce queues at airports. Resulting in benefits to the Aircraft operators, Airports, passengers, and Governments [4].

Labati et al. [5] represented two types of communication performed at an E-Gate: communication between systems interconnected within the E-Gate, and communication with External systems. Internally, an E-Gate is composed of 4 subsystems: Document Authentication System (DAS), Biometric Verification System (BVS), Central Systems Interface (CSI), and Border Guard Maintenance System (BGMS). Externally, the E-gate is part of a larger border control infrastructure. It communicates with external systems to query for additional information to verify the traveler's eligibility to be granted access to cross the border. An E-Gate can query the databases of the following external systems: Visa Management System (VMS), Registered Traveler Program (RTP), and Entry-Exit Management System (EEMS). Figure 3.1 from [5] presents an illustrative overview of the E-Gate logical architecture. Essentially, the data transfer between the above systems, internally and externally, must be secure as

---

[1] *2017 Passenger Summary - Annual Traffic Data.* Jan. 2019. URL: https://aci.aero/data-centre/annual-traffic-data/passengers/2017-passenger-summary-annual-traffic-data/.

[2] Boeing. *Commercial Market Outlet 2020 – 2039.* Tech. rep. 2020. URL: https://www.boeing.com/commercial/market/commercial-market-outlook/.

[3] ReadID. *Which countries have ePassports?* 2020. URL: https://www.readid.com/blog/countries-epassports (visited on 08/07/2021).

it includes personal data, bio-metrics, and decision making information that affect a state's national security.
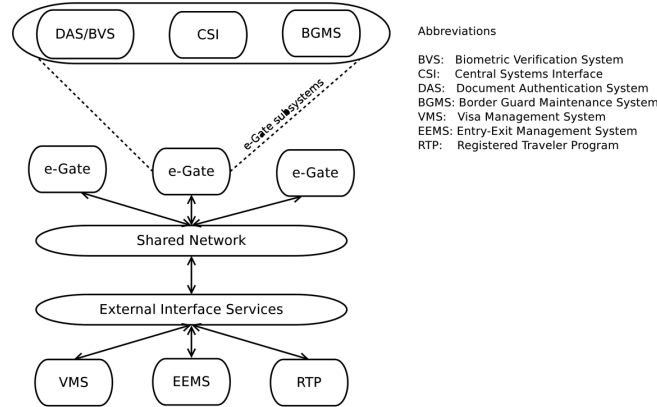
Figure 3.1: Logical Architecture overview from [5]

The workflow of the ABC systems is similar to the one depicted in Figure 3.2. At first, a citizen starts by scanning his passport data page through the gate's scanner. The scanner communicates the image to the Border control system (BCS) which in turn runs its optical data recognition software to extract the data from the scanned data page and verify the document security features. Next, the E-Gate proceeds by reading the embedded electronic chip in the passport. Due to the sensitivity of the bio-metric data, they are not transferred. However, the non-sensitive data stored on the chip, which is equivalent to the data already scanned, is sent to the BCS to cross-check the passport data page. The citizen's facial features and bio-metrics are scanned by the E-Gate to authenticate the citizen against the chip's data. Finally, If the BCS verifies the passport, the E-Gate verifies the bio-metric features, and no match was found in the BCS's watchlists, then the citizen is allowed to pass the border checkpoint.
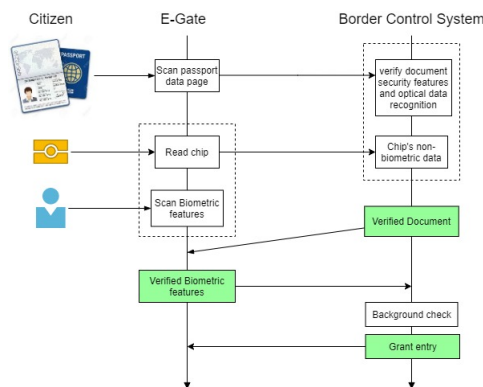
Figure 3.2: Workflow of ABC.

3.2 IoT

IoT is relatively new concept that already has applications in many domains, creating new ways of interactions between humans and small devices, referred to as 'Smart' devices. Applications like 'Smart Airports', Transportation, Home Automation, and many more, are being revolutionized by this technology [6]. IoT is the deployment of network-connected embedded devices, usually constrained, in the physical environment to link it to the digital world through sensors and actuators. IoT devices are classified into several classes depending on their degree of computing resources and power usage. We may safely assume that they may not be able to process sophisticated or even conventional cryptographic operations. In many use cases, the IoT devices handle critical and private data. This imposes the requirement for data integrity and authenticity guarantees, and -because of privacy- in many cases also confidentiality.

Transport Layer Security (TLS) [7], which relies on Transmission Control Protocol (TCP), has been criticized as inappropriate for IoT [8]. Among the reasons are, the infeasibility of IoT devices to maintain long-lived connections due to energy constraints, high header overhead, and the low-latency requirement that is opposed by the delay due to TCP handshake, especially in a lossy network. Alternatively, Datagram Transport Layer Security (DTLS) [9] provides security for communication channels relying on User Datagram Protocol (UDP). In contrast to TCP, UDP is more suitable for IoT. It is an unreliable protocol as it does not care about message delivery , resulting in a lower header overhead. In addition, it introduces less traffic to the network.

The Internet Engineering Task Force (IETF) introduced Constrained Application Protocol (CoAP) [10] that is intended to be a generic application protocol for constrained environments. Similar to the ubiquitous HTTP [11], CoAP realizes a subset of the Representational state transfer (REST) architecture. Therefore, it easily translates to HTTP. Moreover, CoAP leverages UDP as its transport layer protocol which is secured by DTLS. Hence, CoAP is one of well suited protocols for IoT.

3.3 V2V

- define v2v - architectural overview of v2v - use case messages - secure communication (Tesal, etc...)

4 Secure Bootstrapping

The term "Bootstrapping" is used in a spectrum of contexts including Computing, Law, Finance[4]. It is inspired from the late idiom "to pull oneself up by one's bootstraps" which means to succeed or elevate yourself without any outside help. In the context of Networking, bootstrapping is an initial procedure between an unconfigured devices which intend to communicate with a network for the first time. Its goal is to provide the device with the required information that enables it to establish subsequent secure communication channels with the desired network. Such information can be certificates, configurations, and metadata.

Integrity and confidentiality of information flow between the two ends of the communication are are what defines end to end secure channels. Authentication, whether unilateral or mutual, is another fundamental aspect to achieve channel security. Existing protocols, such as TLS, can achieve End-to-end security between parties. However, analogous protocols rely digital certificates and credentials which are managed by local or third party Public Key Infrastructure (PKI). Therefore to ensure correct and secure execution of protocols certificates must be securely provisioned to their corresponding identities. Typically at the manufacturing phase, the manufacturer usually installs globally unique manufacturer provided identifiers known as the Initial Device Identifier (IDevID) [13]. Its main use is for identity verification purposes and it should not be used to enforce data integrity nor confidentiality. Upon successful completion of the bootstrapping process, the device should posses identifiers that allows for subsequent establishment of secure channels with the network domain. An identifier in this set of identifiers is known as Locally Significant Device Identifier (LDevID) [13].

Bootstrapping approaches differ in the degree of manual user involvement and the amount of information on the device which must be pre-configured by the manufacturer. The survey by Sethi et al. [14] provides a classification for Bootstrapping mechanisms into general categories: Managed methods, Peer-to-Peer (P2P) and Ad-hoc methods, Opportunistic methods, and Hybrid methods.

- *Managed methods*: These bootstrapping approaches count on pre-established credentials and trust anchors for authentication and security. The required initial information and cryptographic material can be acquired either at manufacture time in the factory, or through Out-Of-Band (OOB) means, e.g. using a smart card or a USB. Examples of this method are Extensible Authentication Protocol TLS (EAP-TLS) [15] and Generic Bootstrapping Architecture (GBA) [16].

- *P2P methods*: Contrary to managed methods, these bootstrapping methods do not rely on any pre-established cryptographic material or information. Instead, the bootstrapping protocol results in credentials being established for subsequent secure communication. Typically, resulting credentials are authenticated using an OOB channel. This

---

[4]*Bootstrapping*. URL: `https://en.wikipedia.org/wiki/Bootstrapping` (visited on 08/08/2021).

category of methods may be utilized if the manufacturer is incapable or untrustworthy to generate the desired credential.

- *Opportunistic methods*: Unlike previous methods where the authenticity of the initially presented identity is verified, approaches which fall under this category rely on verification of the continuity of the initial identity provided. Bergmann et al. [17] have developed a secure bootstrapping mechanism that is an example of this category.

- *Hybrid methods*: A wide range of deployed approaches use components from both managed and P2P methods. Such approaches are categorized as hybrid methods.

This chapter discusses two voucher-based bootstrapping protocols that aim at being zero touch protocols by eliminating the user interference. These protocols are Bootstrapping Remote Secure Key Infrastructure (BRSKI) [18] and Secure Zero Touch Provisioning (SZTP) [19]. Both protocols fall under the managed methods category.

4.1 Bootstrapping Remote Secure Key Infrastructure (BRSKI)

BRSKI is a product of the ANIMA working group of the IETF. It is an automated bootstrapping protocol that enables an unconfigured device to discover and securely join an unfamiliar network domain it is installed in. BRSKI results in the device device acquiring an X.509 root certificate to authenticate the network domains' elements and establish consecutive secure channels. Moreover, the device can use the obtained certificate to perform further certificate enrollment protocols, like Enrollment over Secure Transport (EST) [20]. BRSKI is capable of realizing a large scale of thousands of devices in a risk prone environment. For example, customer devices provided by ISPs which are directly shipped to the customers.

4.1.1 Architecture Overview

The environment of BRSKI is composed of three general entities: the network domain, the pledge, and the manufacturer services. Figure 4.1 shows an overview of the protocols architecture.
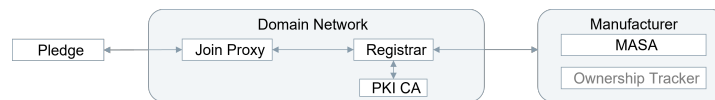
Figure 4.1: BRSKI Architecture.

The network domain is the domain of the alleged new owner of the device, i.e the network that is expecting the device to be connected to it. The domain is a network of entities who share a common local trust anchor. It incorporates a PKI to govern the issuance of digital certificates to provide unique digital identifiers for clients and establish end-to-end security.

A join proxy is another component of the domain. It helps with the discovery of pledges that intend to join the network. In addition, it is responsible for discovering the domain

registrar(s) and determining the proxy mechanisms supported by the registrar and utilizing the lowest impact mechanism. Pledge discovery methods can be classified into passive and active methods. GRASP flooding [21] is a pledge passive discovery method for autonomic networks [22]. It is short for GeneRic Autonomic Signaling Protocol. It is used for signaling between autonomic service agents. GRASP provides discovery, flooding, synchronization, and negotiation functionalities for the technical objectives through respective GRASP messages. Pledge discovery via GRASP multicast flooding is the normative and mandatory method for BRSKI. On the other hand, DNS-based Service Discovery [23] over Multicast DNS [24] as well as DHCP [25] are pledge active discovery methods. They can be used as secondary discovery methods in parallel to GRASP. Moreover, A proxy provides HTTPS connectivity and forwards messages without examination between a pledge and a registrar in the network, and without interfering with the protocol messages.

A pledge is an unconfigured device attempting to join the network domain. Its goal is to be securely bootstrapped in a zero-touch fashion. To achieve this goal, the pledge establishes a TLS connection with one or more of the domain's registrars through the domain's proxy. It is necessary for the pledge and the registrar to establish mutual authentication. A manufacturer installed IDevID is used for pledge authentication to the domain's registrar. It is installed during the manufacturing process and includes certificates signed by the manufacturer and unique identifiers that represent the pledge, in addition to, the pledge unique serial number given by the manufacturer. It is recommended that The provided certificates are used for authentication with the registrar and the signing of voucher requests. The unique serial number is used in vouchers and voucher requests to ensure linkability.

A registrar is an element of the domain that is responsible to carry out the bootstrap process for the pledge. Also, it can be considered as the Registration Authority (RA) for the domain's PKI. A domain can have one or more registrars which all have to be recognized by the domain proxy. If the pledge is capable to concurrently connect to multiple registrars, it is advisable to do so as this protects against a malicious proxy attempting a Denial of Service (DoS) attack like Slowloris.

A manufacturer is the entity that produced the device and set up its initial configuration (IDevID). It provides two distinct services: the Manufacturer Authorized Signing Authority (MASA) service and ownership tracking and validation. MASA can be a third party service that signs the vouchers issued for the bootstrapping process. It is also responsible for providing a repository for audit-log information of bootstrapping events. The service is contacted each time a pledge performs a zero-touch bootstrap in an attempt to enroll into a domain. It takes the decision whether or not issue the voucher according to the MASA policy. Voucher issuance could be done blindly at the lowest security level or it could be tightly bound to the sales channel that verifies the actual ownership of the domain. Hence, the manufacturer can provide protection against stolen devices or illegitimate resale of devices by declining voucher issuance to the suspected pledge.

Ownership tracking and validation is an optional manufacturer service. It is supposed to log all claim attempts and to know which device is owned by which domain and provide such information to registrars. A verified log entry indicates that the pledge was issued a voucher as a result of positive verification of ownership.

TODO formal verification

4.1.2 Protocol Details

This section describes the message sequence of BRSKI illustrated in figure 4.2 and elaborates on content of the exchanged messages. The numbering sequence referenced through this section refers to the message numbers in figure 4.2.
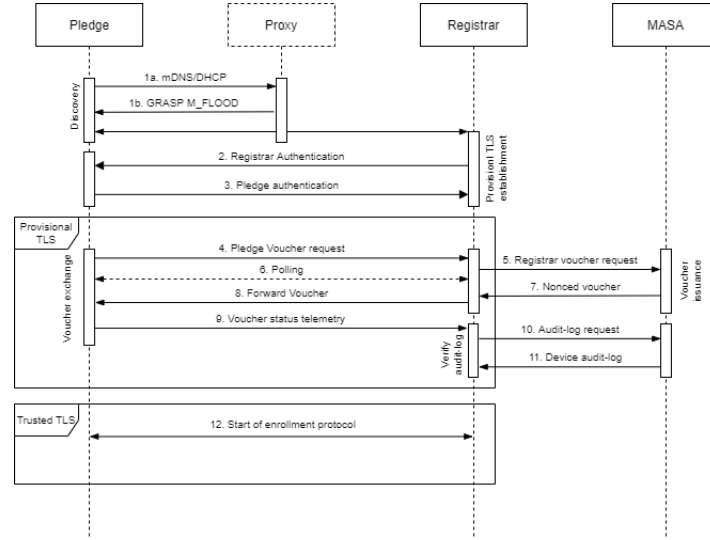
Figure 4.2: A successful BRSKI protocol run.

- *Message 1a, 1b (Discovery phase):* It is the first phase of the protocol where the pledge identifies the domain proxy. This can be performed by a pledge initiated mechanism as in message (1a) or via a proxy initiated mechanism as in message (1b). After successful discovery, the pledge can address a domain registrar through the proxy. A proxy does not assume any specific TLS version.

- *Message 2, 3 (Provisional TLS establishment):* The pledge attempts to establish a TLS channel with each discovered registrar to ensure End-to-End security. The pledge must not use any TLS version lower than TLS 1.2, while TLS 1.3 is the encouraged version to be used. To establish the channel, mutual authentication has to be performed. At first, in message (2), the pledge receives the registrar's server certificate. However, the pledge does not posses any trust anchors to verify it yet. Therefore, the pledge accepts the registrar's certificate provisionally. Next in message (3), the pledge is authenticated via the installed IDevID. The registrar must be able to verify the provided certificate, however the distribution of the trust anchors for this task is out-of-scope of BRSKI. Meaning, the information received by the pledge must be untrusted, although it is in a TLS channel, till a trusted trust anchor to verify the certificate is received.

- *Message 4:* Having established a secure provisional TLS channel, the pledge initiates the voucher exchange by sending a pledge voucher request to the registrar in message (4). The request must contain a unique nonce per bootstrapping attempt to protect against replay attacks. Also, the request contains the 'proximity-registrar-cert' and the pledge serial number. The 'proximity-registrar-cert' is the End Entity (EE) certificate of the registrar which is also used to establish the provisional TLS channel. Pledge

serial number is a manufacturer defined unique identifier for each device. It is different from the IDevID certificate serial number. Since not all devices have a real-time clock, depending on the device capabilities, the request is recommended to have the 'created-on' value. Finally, the request must be signed using the pledge's IDevID certificate.

The registrar authorizes the pledge based on the authenticated information presented in the pledge's IDevID and the registrar's policy. The policy can be either to allow any device from a specific vendor, to allow any device of a specific type, or to allow a specific type of devices from a specific vendor.

- *BRSKI-MASA TLS channel:* The registrar initiates a TLS 1.2 or newer channel with the MASA where all subsequent communication between the two parties occur within this secure channel. The MASA URL is obtained from the pledge IDevID, as mentioned earlier. To authenticate the MASA, the registrar should be configurable with trust anchors on a per vendor MASA basis as part of the sales process. Moreover, the registrar should also support client authentication mechanisms such as TLS client certificate, HTTP Basic, Digest, or Salted Challenge Response Authentication Mechanism (SCRAM); however TLS Client Certificate based authentication is the recommended method.

- *Message 5:* After obtaining the pledge's voucher request, the registrar constructs a registrar voucher request that is sent to the MASA to obtain a voucher for the pledge. The registrar voucher request is a JSON document that is signed using a CMS structure. The JSON document encapsulates the pledge voucher request CMS object that was sent to the registrar and is referred to as 'prior-signed-voucher-request'. Moreover, the request contains the 'created-on' field which holds the timestamp the request was formed on. In addition, it consists of other fields which relate to the pledge request like the pledge serial number, the nonce used produced by the pledge and used in the pledge request, and 'idevid-issuer' field which holds the issuer value of the pledge IDevID certificate. The registrar includes some certificates in the registrar voucher request CMS object as well. Those certificates are used by the MASA to be pinned into the voucher to be later used by the pledge as a trust anchor for authenticating the domain registrar. Therefore, the certificates enclosed by the registrar in the request have to be part of the chain it wishes the MASA to pin in the voucher. Hence the specificity of the attached certificates is considerably significant. A 'pinned-domain-cert' can be as specific as the registrar's TLS EE certificate. On the other hand, if it is as general as a public webPKI Certification Authority (CA) it could permit any entity that possess a certificate issued by that authority to claim ownership of the device.

On the other hand, the pledge might not be available at the time of deployment to send a pledge voucher request, or the registrar speculates to not being able to reach the MASA at the time of deployment where the pledge will be available. Such use cases justify the need for nonceless registrar voucher request. In these cases, the previous message (4) would not exist. To formulate this request, the registrar has to acquire the pledge's serial number and IDevID issuer, however, they are obtained through out-of-band means. Subsequently, the nonce field of the request is omitted.

- *Phase 6 (Polling)*: Before processing the pledge's request, the registrar may send the pledge an HTTP 202 response message which indicates that the request received earlier

has been accepted for processing however processing is not yet complete. This response initiates a polling phase between the pledge and the registrar. A "Retry-After" field is specified within the headers of the registrar's response that indicates the minimum time for the pledge to wait before asking for a response for the voucher request sent earlier. After the specified waiting time, the pledge polls the response by resending the exact same request and must not change the nonce nor sign a new voucher request. If the pledge is simultaneously trying to bootstrap itself with several registrars of the network, it can be overwhelming for the pledge to keep track of all the "Retry-After" times. Therefore, a pledge may ignore the specified interval and follow a hard-coded "Retry-After" interval. A pledge should be able to hold the retry state for a maximum of 4 days.

- *Message 7:* upon receiving the voucher request, the MASA performs a set of checks to decide weather to issue the requested voucher. Given the fact that vouchers have a short lifetime, the request may be from a registrar that has been issued a voucher previously, i.e a voucher renewal request. In this case, the request should be automatically authorized by the MASA.

  The MASA extracts the certificate chain attached in the signed CMS object. If the domain CA is unknown to the MASA it is considered as a temporary trust anchor as the intention is not to authenticate the message rather to establish consistency of the domain PKI. According to the MASA's policy, it decides which certificate of the chain supplied by the registrar it chooses to pin. It may be the farthest certificate of the chain, or it may be as close as the EE TLS certificate of the registrar. If revocation information is available for that certificate, it must be checked by the MASA to prevent issuance of new or renewed vouchers to unauthorized registrars. Next, the CMS signature is validated using the domain's CA extracted from the voucher request. Also, the signing certificate is verified to contain the 'id-kp-cmcRA' Extended Key Usage. This ensures that the signer is an entity that is authorized to be a registrar of the domain. Hence, assures domains that a MASA only accepts requests from domain registrars.

  In case of nonceless requests, It is mandatory for the MASA to authenticate the registrar. The decision to issue a noncless voucher is taken according to the MASA policy that is out of scope.

  In case of nonced voucher requests, the MASA verifies that the 'prior-signed-voucher-request', enclosed in the registrar request, contains a 'proximity-registrar-cert' that is coherent to the certificate used to sign the registrar voucher request. Moreover, the nonce is verified to be consistent between the registrar voucher request and the 'prior-signed-voucher-request'.

  Subsequent to a successful validation of the request, the MASA responds with an issued voucher in message (6). Any issued voucher by the MASA is recorded in the audit-log. Otherwise if a problem occurs, a response with the appropriate http signaling as described in [18]. For example, a 403 status code response if the voucher request is not signed correctly, or a 406 status code response if the requested voucher type or algorithms cannot be issued due to the MASA's awareness that such pledge is not capable of processing them.

- *Message 8:* The registrar evaluates the received voucher solely for transparency and

future audit-log verification. The received voucher is forwarded to the pledge without any interference or modification from the registrar.

- *Message 9:* After the pledge successfully receives a voucher, the pledge must indicate its status regarding the voucher to the domain. This occurs by sending a status message to the registrar. The pledge decides weather to accept the voucher or not through the voucher validation process. If acceptable, the message should contain the version of BRSKI and a boolean status field to indicate the acceptance status. In case of an unacceptable voucher or a failure, the pledge is expected to fail gracefully. The message should contain a Reason field with a string commenting on the cause. Nevertheless, the Reason should not be excessively descriptive as it may be sent to an unauthenticated and potentially malicious registrar.

  Bearing the voucher, the pledge verifies its validity. It verifies the signature using the manufacturer installed MASA trust anchor. It verifies also that the serial number enclosed in the voucher matches its own. For nonced vouchers, the pledge verifies the voucher nonce corresponds to the nonce it sent earlier in the voucher request. However nonceless vouchers can be accepted according to pledge local policy. The pledge can be configured to always accept nonceless vouchers to realize the use case where the MASA is unreachable at the time of pledge deployment.

  A pledge could be operating in other similar security reduced mode that skip voucher validation in favor of offline or emergency touch-based deployment bootstrapping procedures. For example, Trust on first use (TOFU) or physical presence methods such as the use of serial console or depressing a physical button during bootstrapping. However, TOFU must not be available unless a hardware-assisted Network Endpoint Assessment (NEA) is supported. Meanwhile, it is only recommended for other methods of skipping voucher validation. This recommendation serves as a prevention against unintended use of offline methods when autonomic methods fail or are unavailable.

  Upon successful verification of the voucher, the voucher's pinned-domain-cert should be considered by the pledge as a trust anchor. The current provisional TLS connection between the pledge and the registrar is evaluated using the obtained trust anchor. The pledge verifies the registrar's TLS server certificate using the trust anchor's public key. If the registrar's credentials could be verified, either by directly matching the server certificate or through verifying a higher certificate in its chain, the pledge trust the TLS connection and it is not considered provisional any further.

- *Message 10:* After receiving the pledge status telemetry message, the registrar requests the MASA audit-log from the MASA. The log data helps the registrar make a knowledgeable decision regarding further proceeding of the bootstrapping process. The decision making criteria is based upon the security requirements of the registrar domain. Hence, the criteria is out of the protocol's scope. The request content is the exact same registrar voucher-request sent earlier to the MASA, but is directed through a different URI specific for requesting the audit log, which is "/.well-known/brski/requestauditlog". Reusing the same message minimizes the required cryptographic and message operations on both ends. The registrar may reuse the cached voucher request and the MASA may take advantage of its internal state to correlate the message with the already verified request averting additional operations.

- *Message 11:* A MASA can infer the proper pledge log to be prepared from the "idevid-issuer" and the "serial-number" information included in the received request of the previous message. Instead of immediately responding with the audit-log, the MASA can a HTTP 201 "Created" response with a URL in the "Location" header field redirecting to actual audit-log. The response log is a JSON format document consisting of all the log entries associated with the pledge. Nevertheless, a MASA that sends out URLS has to ensure they are unpredictable to avoid enumeration attacks against device audit-logs.

  The log format structure consists of several entries: "version", "events", and "truncation". "version" is an integer value representing the log format version. "events" is an array of event objects that are associated to the device. Each of the event objects is comprised of a set of entries. The "date" entry represents the event's timestamp in the format according to [26]. The "domainID" is a unique identifier for the domain's registrar that encodes the pinned-domain-cert's SubjectKeyIdentifier or SPKI fingerprint in base64. A "nonce", if exists, is a base64 encoding of the same nonce used in the voucher request and issued voucher. If it is a nonceless voucher, then the field should preferably be set to null rather than omitting it. The "assertion" field indicates the level of verification with which the MASA issued the voucher. It can have one of three values: "verified", "logged", and "proximity"; the latter being the one supported by this protocol. "truncated" field shows the number of event truncations for the specified domainID. Lastly, since audit-logs can be arbitrarily large, duplicated or old entries may be truncated as an optimization for the log structure. The "truncation" entry contains meta-information about truncated entries such as "nonced duplicates", "nonced duplicates", and "arbitrary".

  On the registrar side, the received audit-log is vetted for discrepancies and unexpected behavior like the pledge previously imprinting to an unexpected domain or whether a certain domain possesses a nonceless voucher and can reset the device anytime. If the registrar's audit-log verification is successful, then the bootstrapping process is complete.

- *Message 12:* At this point, the pledge has a trust anchor allowing it to verify the registrar, as well as a trusted TLS channel between them. Therefore the environment is suitable to start a certificate enrollment protocol after which the pledge obtains digital certificates that authenticate it to the domain and authorize it to utilize the relevant domain services.

  BRSKI is described as an extension for EST that provides automated proposal instead of the originally manual authentication method that relies on the intervention of a human user. Hence it is recommended for a pledge to use EST following BRSKI as a certificate enrollment protocol as it is considered a harmonious integration.

 Nonetheless, the succeeding certificate enrollment protocol is not limited to EST, however, a variety of certificate enrollment protocols can be used. Using EST is an example of a pull model where the EST server is the protocol initiating party. As an example of the push model architecture, Certificate Managment Protocol (CMP) can be used as the certificate enrollment protocol since the EE is the initiating party of the protocol.

4.2 Secure Zero Touch Provisioning (SZTP)

SZTP is as well an outcome of the IETF group. It introduces a bootstrapping technique for devices in initial factory state. It aims at securely supplying networking devices with bootstrapping data without further actions required other than physical placement and connecting the power and network cables. SZTP is primarily concerned with physical devices, however, it has the potential to be extended to logical entities like virtual machines. The protocol is not limited to provisioning trust anchors to the device but rather can update the device's boot, commit an initial configuration, and execute arbitrary scripts to handle auxiliary requirements.

4.2.1 Architecture Overview

The protocol's architecture is divided into three abstract entities: the owner, the manufacturer, and the device. Figure 4.3 illustrates the protocol's architecture.

Figure 4.3: SZTP Architecture.

The owner is used to refer to the person or organization that owns the device. Throughout the protocol, the owner term abstractly represents the owner network domain and its sub-entities without explicitly referring to them. An example of the owner's sub-entities are the domain CA and the domain registrar. An owner possesses an *Owner Certificate* which is an X.509 certificate that identifies the owner's identity and binds it to its public key. In addition to validating the owner's identity, the certificate is used to by other entities to validate the owner's digital signature over received artifacts. A Network Managment System (NMS) is assumed to be part of every owner's structure. A NMS is a software used by network administrators to monitor software and hardware nodes in a network and logs data from those nodes for reporting. The bootstrapping process introduces newly admitted devices to the NMS.

The manufacturer is the entity that produced the device. The term is also used to refer to entities that the manufacturer delegates functions to. A MASA is an example of a third party manufacturer delegate entity. Each manufacturer is supposed to operate a MASA or most commonly delegate its functionality to a third party. The MASA is responsible for generating the voucher required by the device to authenticate the owner and bootstrap the device. Each of the manufacturer services and its related entities has its own certificate for authentication and signing. The manufacturer preserves a list of trusted voucher-signing authorities and well-known bootstrap servers. The shipped devices from the manufacturer are pre-configured with the list of trust anchors along with their own IDevID credentials. Owners are assumed to trust the same pre-configured trust anchors as the ones maintained by the manufacturer.

The protocol defines three artifacts that are exchanged through out the protocol: Conveyed information, ownership voucher, and owner certificate. Those artifacts provide the device

with all the required information needed for bootstrapping. The conveyed information is divided between redirect information and onboarding information. Bootstrap servers are Representational State Transfer Configuration (RESTCONF) servers that provide bootstrapping artifacts to devices. Depending on the type of information a server provides, they could split into two types of servers: Redirect servers and Onboarding servers. Redirect servers only returns redirect information to clients while onboarding servers only return onboarding information. However, servers are not limited to mentioned categories only. There may exist a server that can provide both types on conveyed information.

Redirect information redirects a device to another bootstrapping server. Redirect information encodes a list of bootstrap servers hostnames and an optional trust anchor certificate that the device can use to authenticate each bootstrap server with. Depending on the source, redirect information may be trusted or untrusted. It is trusted whenever obtained via a secure connection to a trusted bootstrap server or whenever it is signed by the device's owners. Trusted redirect information is useful for enabling a device to establish a secure connection to a specified bootstrap server. Untrusted redirect information is useful for directing a device to a bootstrap server where signed data has been staged for it to obtain. Redirection acts as a guide for device to discover the bootstrapping servers capable of providing onboarding information.

Onboarding information is a bundle of data required for a device to perform the bootstrapping process into the owner's network. It includes information about the boot image a device is required to be running, an initial configuration a device must apply, and scripts to address arbitrary needs which the device must successfully execute. Onboarding information must be obtained from a trusted source, either through a secure connection to a trusted bootstrap source or the information obtained is signed by the device's owner.

Devices can obtain the needed bootstrapping data from various sources according to the owner's deployed infrastructure. However, The concern is focused on sources that can supply devices with the information in an automated manner. Here are examples of touchless bootstrapping sources.

- A DNS server can be utilized as a form of a bootstrapping server. It is an attractive method for environments which already employ a DNS infrastructure. It does necessarily require communication with an internet-accessible third party DNS service. Using DNS as bootstrapping server is considered a touchless bootstrapping as it does not involve any external interaction. However, a DNS server is not a trusted source of bootstrapping information, even if DNSSEC [27] is used to authenticate the DNS records. The reason being that the device cannot verify if the returned domain belongs to its rightful owner. Therefore, the returned DNS records must either be signed for later verification by the device or the records must be processed provisionally.

  Devices supporting DNS server as a bootstrap server are offered two prioritized types of queries to the server: device-specific and device independent queries. Queries must first be attempted using multicast DNS before unicast DNS. A DNS response for the device-specific query can encode the three artifacts into the TXT records but the response must be signed. However, by DNS conventions, the size of signed data is large. Therefore, it is implausible for the signed onboarding information to fit the UDP-based DNS packet size. Thus, if signed onboarding is to be sent over DNS, it is expected to be transported over

TCP so as to handle the DNS TXT record size. Otherwise, only redirect information shall be returned. On the other hand, DNS response for device-independent queries does not support returning onboarding information because the DNS server is incapable of returning signed data in response to the device-independent query. Accordingly, only redirect information shall be returned.

- Similar to DNS servers, a DHCP server is another untrusted bootstrapping data source. It can be leveraged by deployments that already employ a DHCP infrastructure. However, a DHCP server is a limited bootstrapping source due to its lack of ability to transmit enough data to hold signed bootstrapping data. Nevertheless, a DHCP server can return to a device unsigned redirect information.

- Lastly, a specific bootstrap server can be deployed to provide bootstrapping data and receive data from devices. It is defined as a RESTCONF server that implements the YANG module described in [19, Section 7]. Moreover, it may be using TLS which may eliminate the requirement to sign the transmitted bootstrapping data if the bootstrap server is trusted. Otherwise, if the bootstrap server is not trusted by the device, the conveyed information must be signed or processed provisionally. Whether a bootstrap server is trusted or not depends on the device's knowledge of the server's trust anchor. A bootstrap server exposes two endpoints to communicate with devices. Namely, the "get-bootstrapping-data" for devices to request bootstrapping data and the "report-progress" for devices to report their bootstrapping process status back to the bootstrap server, such as warnings, errors, and the result of the process. A bootstrap server may be hosted by the owner or is an internet accessible server hosted by the manufacturer.

4.2.2 Protocol Details

SZTP is focused on the enrollment process between the device and its owner. However, before initiating SZTP, the owner and the manufacturer must complete a pre-protocol phase where the owner orders the devices and enrolls itself to the manufacturer where they exchange vital information required for the owner to be able to run SZTP. This phase can be referred to as the Device Ordering and Owner Enrollment phase. At first, the owner orders the needed devices from the manufacturer. In turn, the manufacturer always provides the owner with the trust anchor it will need to verify the IDevID certificates the devices use. If the manufacturer offers an internet-based bootstrap server, the manufacturer may also provide the owner with the necessary credentials to access the hosted bootstrap server to configure it. Consequently, the owner configures its NMS with the information obtained from manufacturer. If a remote bootstrap server credentials were obtained, the NMS configures an owner trust anchor certificate onto the bootstrap server that is to be used as the 'pinned-domain-cert' of the voucher at the time of dynamically generating the enrollment voucher. After the manufacturer concludes the order and the devices are shipped, it may send the owner the serial numbers of the devices, other meta information, or even nonceless ownership vouchers. The owner may configure its NMS with the received information to prepare its network infrastructure to enroll the expected devices.

Eventually, the owner is in possession of the devices and physically sets them up in its environment. As soon as the unconfigured device is booted for the first time, it checks if its factory default configuration enables SZTP bootstrapping. If enabled, the device initiates

the SZTP protocol and goes through a series of phases to bootstrap into the owner network infrastructure. Figure 4.4 mentions the said phases which are discussed below.

Figure 4.4: SZTP protocol phases.

a) *Discovery:* First of all, the device discovers and enumerates all the available sources of bootstrapping information. A legitimate discoverable source of bootstrapping data can be one of those discussed in section 4.2.1. For each source supported by the device, the device attempts to obtain bootstrapping data from it. The Bootstrap sources approached in order of their close proximity to the device. A device only process one bootstrapping server at a time.

b) *Authentication:* Secondly, when contacting a bootstrap server, it is mandatory for the device to authenticate itself to the server. The device authenticates itself using its TLS client certificate which is the same as its IDevID certificate. The server validates the client TLS certificate through the manufacturer trust anchor obtain earlier. Moreover, the device may possess more trust anchor certificates other than its TLS client certificate. If the device supports connecting to remote well-known servers, then it must authenticate the server through a pre-configured list of trust anchors for well-known servers. In addition, the device requires trust anchor certificates from the manufacturer for validating the ownership voucher. Devices should have the certificates chain of trust anchor certificates up to and including the self-signed root certificate.

c) *Bootstrap data transmission & processing:* After establishing a connection to the bootstrap server, whether trusted or not, the device starts receiving artifacts to bootstrap itself. The type data acceptable from bootstrapping sources depends on their trust state from the devices point of view. Untrusted sources can provide signed or unsigned redirect information, bearing in mind that the server which the device was redirected to must provide the device with further redirect information. Contradictory to redirect information, untrusted sources can only provide signed onboarding information. On the other hand, trusted sources can provide both redirect and onboarding information regardless signed or not. Since trusted sources already establish an authentic and secure channel with devices, it is redundant to have the transmitted information signed, therefore unnecessary. Naturally, the ownership voucher and the owner certificate are signed in any case and relevant revocation information may be additionally included in case it cannot be obtained dynamically. If the device received bootstrapping information that does not obey the restrictions respective to its source's trust status, the device should discard the information and quit the bootstrapping process with the said source.

Whenever a device receives signed data, it must validate the data signature before processing it regardless of the data type. Signed data must always be accompanied by an owner voucher and an owner certificate. To validate the signed data the device must first validate the owner voucher and certificate. The device begins by validating

the voucher's signature using its pre-configured trust anchors. It also verifies that the voucher's creation timestamp was in the past if the device is equipped with an accurate clock, and that the voucher "assertion" value is acceptable by the device. In addition, the serial number of the device has to match that of the voucher as well as the "idevid-issuer", if present. If valid, the device extracts the now trusted "pinned-domain-cert" certificate from the voucher. Next, the device authenticates the owner certificate by verifying its certificate path to the extracted "pinned-domain-cert", and if specified in the voucher, the revocation status of the certificate chain used to sign the owner certificate has to be checked. Finally, if the owner certificate is validated and the conveyed information was truly signed by that certificate, only then can the signed redirect or onboarding information be processed.

Redirect information provides a device with a list of bootstrapping servers. The device proceeds through the list till it reaches a bootstrap server it can bootstrap from. If a server provides the device with bootstrapping data, the device must attempt to process it before proceeding to the next bootstrap server. A bootstrap server may further redirect the device to other bootstrap servers. In this case, implementations must limit the number of recursive redirection to a maximum of ten redirects to avoid indefinite redirections. Trusted redirect information may be accompanied by a trust anchor certificate. Hence, the device must authenticate the redirected to server certificate against the certificate obtained from the preceding bootstrap server. If no trust anchor is provided or the redirect information is untrusted, the device must process the information provisionally.

Onboarding information is parsed by the device to extract the required instructions to be executed to evidently bootstrap the device. Onboarding information defines the boot image a device must run, pre-configuration scripts, initial configuration to be committed, and post-configuration scripts. The device must process the received information in that order. The boot image is verified if it complies with the required boot image. If the requirement is not satisfied, the device must download the boot image from the URIs defined in the onboarding information. After verifying the installed image, the device must reboot which will lead it to restart the bootstrapping process but with the new boot image, but then the device will not halt at this step and will proceed to further steps. If pre-configuration scripts are specified, the device must execute them and log their outputs. If an initial configuration is specified, the device must apply it according to the approach stated in the "configuration-handling" node. Similar to pre-configuration scripts, post-configuration scripts must be executed and their output must be logged.

Whether a device should return progress reports of the bootstrapping process is decided according to the trust state of the bootstrapping server. A device must not send progress reports to untrusted sources. However if the source is trusted, the source devices the minimum verbosity level it requires from a device in the "reporting-level" node of the "get-bootstrapping-data" RPC-reply. Moreover, the device may send more reports than originally specified by the trusted source. If an error occurs during any of the stages, devices must roll back the current step and previous steps and exit the bootstrapping process. However, results of some steps of the bootstrapping process may be retained, like an updated boot image. When eligible, devices must send error reports whenever they occur at any of the onboarding information processing stages. The bootstrapping

process with current source should be quitted and the device should start the process over with the next bootstrapping source, if available.

4.3 Protocols Comparison

BRSKI and SZTP are two different protocols which tackle the same problem. Their main goal is to automatically and securely provision trust anchors to devices without any manual interference at the device. Hence, allowing such unconfigured devices to authenticate the identity of their new owners and establish secure channels for further processes like certificate enrollment. This section compares and contrasts both protocols. In addition, Appendix A provides a comparison between the two protocols and their terminology in a tabular form.

As mentioned, BRSKI and SZTP share a similar aim. BRSKI aims at storing a root certificate on the pledge that is sufficient for verifying the owner's domain registrar identity. While SZTP is more comprehensive as it involves updating the pledge boot image, commit an initial configuration, and execute arbitrary scripts. The fundamental architecture for both protocols is the same as both function on three elemental entities: the pledge, an owner representative known as registrar, and the manufacturer's MASA. However, unlike BRSKI, the communication between the owner's registrar and the MASA is not intrinsic to the protocol. BRSKI is integrated with EST as an enrollment protocol to be utilized after bootstrapping, but it is not limited to it. On the other hand, SZTP is not natively integrated with an enrollment protocol. Both protocols use HTTP on top of TLS as their transport protocol, but BRSKI specifically restricts the use of TLS to version 1.2 or higher. In addition, BRSKI proposes the usability of CoAP as an alternative for HTTP. Both protocols do not require any specific cryptographic algorithm.

The bootstrapping data differ between both protocols. For SZTP, the bootstrapping data are composed of the types of conveyed information explained earlier, the ownership voucher, and owner certificate. However, for BRSKI, it is only the ownership voucher artifact. Therefore, both protocols are reliant on the voucher artifact to provide the trust anchor necessary to verify the owner certificate. The owner requests the voucher similarly in both protocols. Nonceless vouchers are provided to owners out-of-band, while nonced vouchers are requested dynamically during a protocol run. In general, bootstrapping data are protected. In BRSKI and SZTP's untrusted channels the data must be signed, but for SZTP's trusted TLS channels the data may be optionally signed.

Owners supporting either protocol provide means to discover pledges and for pledges to obtain bootstrapping data from. For BRSKI, it depends on a domain proxy to discover pledges mDNS or GRASP broadcast messages. It introduces only one bootstrap source which the domain registrar. The registrar acts in principal as the owner's CA and is responsible for bootstrapping pledges as well as communicating with the manufacturer services. Meanwhile, SZTP provides more options for pledges for discovery. An owner can administer a DNS server, a DHCP server, or a redirect server to respond to pledges discovery broadcast message and redirect the pledge to a source it can bootstrap from. However, similar to BRSKI, pledges can only obtain bootstrapping information from one source, namely, a bootstrap server. That is a server which is capable of providing redirect information as well as onboarding information.

Nevertheless, bootstrapping sources of both protocols can be hosted locally at the owner side or remotely by a trusted third party.

Both protocols allow owners to authorize the pledges which are allowed to initiate the protocol according to their vendor-specific serial numbers which are known in advance to a protocol run. Moreover, BRSKI can optionally couple the serial number with a specific pledge type and/or a specific vendor.

An owner needs to know the URI for the MASA responsible for issuing vouchers for the pledge running the bootstrapping protocol. BRSKI recommends that the IDevID of the pledge contains URI for MASA to contact. On the other side, SZTP does not propose an inherent method to obtain the MASA URI. It rather relies on the OOB owner-manufacturer enrollment process to acquire the URI. After contacting the MASA, owners decide on whether to accept to bootstrap the pledge. In both protocols, the decision is made according to the verification of the attributes of the issued voucher. Also, BRSKI checks the MASA audit-log for further verification. As SZTP is more concerned with the pledge-owner communication, checking of the MASA audit-log is not discussed. However, SZTP does not limit it.

BRSKI allows for pledge ownership transfer through issuance of new vouchers. Even though it is also feasible in SZTP, it is out of the protocol's scope. On the contrary, SZTP allows owners to provision domain specific configuration to pledges during the bootstrapping process, but BRSKI is not capable of providing such information.

A pledge is the protocol initiator for both protocols. Pledges compatible with either protocol must have an initial pre-configured state with a set of minimum required parameters. The initial state must contain the pledge IDevID and the trust anchors needed to verify the MASA. Optionally for both protocols, the device can be configured with a list of well-known remote bootstrap servers and their respective trust anchors. Moreover, SZTP pledges may be configured with a TLS client certificate and its related intermediate certificates. The initial state trust anchors are not updatable via either protocol methods.

The bootstrap source must be authenticated even if it is not yet trusted. For either protocol, pledges establish a TLS channel with the bootstrap source. But if the pledge is not able to validate the source against one of its trust anchors, the TLS channel is provisionally trusted.

For both protocols, pledges must check the validity of received vouchers with respect to their expiry time. This is dependent on whether the device has an accurate clock and on whether the voucher is nonced or nonceless, as expiry time might be omitted in nonced voucher. Pledges must also check the revocation status of received certificates, if revocation information is provided or if revocation checks are enforced.

On the Manufacturer side, BRSKI optionally supports manufacturer tracking, while SZTP does not. Each protocol relies on the MASA for voucher issuance. However, BRSKI relies on the MASA for further functionalities it supports such as voucher renewal and maintaining and providing the voucher audit-log.

5 Post-Compromise Security

Intro
compromise window

5.1 Extended Triple Diffie-Hellmann (X3DH)

FIXME intro

Key establishment protocols are utilized by communicating parties to establish shared secrets, in some cases, with the aid of a trusted third party [28]. As per [28], key agreement is a key establishment technique in which a shared secret is derived by parties as a function of information related to them, such that the secret in not predeterminable nor deducable by outsiders.

The Security of protocols is required to be proven to avoid the devastating effects of malicious attacks. Therefore, verification of security protocols is a vital process. Verification is proving that a protocol model is secure by achieving a set of security goals. There are various model checking tools that provide verification.

X3DH is a key agreement protocol intended for asynchronous settings [29]. The protocol can establish a shared secret between two parties in an environment where it is recurring that one of the parties is offline and the other wishes to send it an encrypted message. Also, X3DH provides the desired feature of forward secrecy.

This chapter discusses the Extended Triple Diffie-Hellmann (X3DH) key agreement protocol [29] and verifies the security of the protocol using the OFMC model checker [30].

5.1.1 Protocol Overview

5.1.1.1 Roles

A protocol run involves three roles: Alice, Bob, and a server. In this description, Alice wants to send an encrypted message to Bob. Bob is the party that may be offline at that time and wishes to enable other parties to derive a shared secret with it, through a set of public information Bob publishes. The server is responsible for 1. storing the public information published by Bob, and 2. storing messages for offline parties till they are fetched. The server is not trusted, however, it is assumed to be resilient against DoS.

Owning Role	Notation	Description
Alice	IK_A	Long-term identity Key
	EK_A	Ephemeral Key
Bob	IK_B	Identity Key
	SPK_B	Signed Prekey
	OPK_B	one-time Prekey

Table 5.1: X3DH keys.

5.1.1.2 Keys

X3DH utilizes Elliptic Curve asymmetric key pairs. All keys used in a protocol run must all be derived from the same curve, either X25519 or X448. Each role has to have a set of keys to run the protocol. Table 5.1 lists the public keys required for each role. Note that for each public key, there exists a corresponding private key at its owner.

- *Identity keys*: They are long-term public keys known by their corresponding parties before the protocol run.

- *Ephemeral Keys*: This type of key pair is freshly generated within the protocol run.

- *Signed Prekey*: This key pair is generated and signed by its owner. The prekey is signed using the private identity key. The life time of this key pair is shorter than that of the identity key pairs as it is updated periodically by its owner. The corresponding role owns only one signed prekey at a time.

- *One-time Prekey*: The corresponding role generates multiple one-time prekey. Each is can be used for only one protocol run. The responsible party is supposed to supply the these keys as they should not run out. In case there are not any keys left, the protocol can run, however without one of the DH operations as explained in section 2.3.

5.1.1.3 A Protocol Run

At first, the protocol starts with a registration phase. Initially, The party acting in the role of Bob publishes to server the public information required to run the protocol with it by any party acting as Alice. Bob publishes his IK_B, SPK_B, Bob's prekey signature, and a set of OPK_B.

Next, the inline party sends the initiates the protocol by sending the first message. First of all, Alice fetches a *prekey bundle* from the server to contact Bob. This bundle contains:

- Bob's Identity key IK_B.

- Bob's signed prekey SPK_B.

- Bob's prekey signature.

- If exists, a one-time prekey OPK_B. The server deletes the OPK_B sent to Alice.

Before proceeding, Alice verifies the prekey signature and quits the protocol if the verification fails. At this point, Alice has enough information from Bob to deduce a shared secret. Alice

generates her Ephemeral key pair EK_A. With the set of available keys, Alice performs three DH operations which are extended to four if a OPK_B is available. Figure 5.1 presents the relation between the keys.

- $DH1 = DH(IK_{A_p}{}^5, SPK_B)$

- $DH2 = DH(EK_{A_p}, IK_B)$

- $DH3 = DH(EK_{A_p}, SPK_B)$

- $DH4 = DH(EK_{A_p}, OPK_B)$

Figure 5.1: X3DH operations [29].

The DH outputs are fed into a KDF to generate the shared secret SK. Next, Alice deletes her EK_{A_p} and all DH outputs for forward secrecy. At this moment, Alice is ready to send the initial message with the content encrypted using SK. The initial message from Alice to Bob has to have enough information for Bob to generate the same SK. The initial message contains:

- IK_A

- EK_A

- Identifiers of Bob's prekeys used by Alice

- An initial ciphertext. The ciphertext can be used as an initial message for a post-X3DH communication protocol, e.g. Double Ratchet algorithm.

Upon receiving the initial message, Bob attempts to derive the SK. Using the key identifiers sent by Alice, Bob loads the private keys corresponding to the public keys Alice used. In combination with the keys IK_A and EK_A Alice sent, Bob can compute the same SK by doing the three (or four) DH operations.

Next, Bob attempts to decrypt the ciphertext. If the message is successfully decrypted to the expected format, e.g. the format of the first message of the post-X3DH protocol, then the protocol run was successful. Otherwise, Bob aborts the protocol and discards the SK.

---

[5]p indicates the private key of the key pair.

5.1.1.4 Security Considerations

Authentication is essential for both parties to guarantee the identity of who they are communicating with. Thus, Alice and Bob must authenticate the keys IK_A and IK_B. However, the protocol specification does not discuss authentication methods.

The one-time prekey used in the fourth and optional DH calculation is for protection against replay attacks as they ensure freshness of the protocol run. Absence of a one-time prekey could lead a replayed message to be accepted by Bob believing that Alice had sent it in the current protocol run.

A server can be a cause of attacks if malicious. It can carry out a Denial of Service attack if it refuses to forward the messages. It can deliberately not distribute one-time prekeys, exposing the protocol to replay attacks. Also, one party can drain all the one-time prekeys, if the server is not attentive to such action, leading to replay attacks.

5.1.2 OFMC Modeling

FIXME add info about the goal of ofmc modelling and the type of adversary tested against

Presented in listing 5.1 the code for modeling the protocol in OFMC using the *AnB* language. This section explains the code of the modeling.

```
1  Protocol : X3DH #Protocol name

3  Types :
4         Agent A, B, s;
5         Function kdf, pk, ik;
6         Number g, NA, EKA, OTP, PREKEY, MSG1, MSG2;

8  Knowledge :
9             A : A, B, s, g, pk(s), pk(B), pk(A), inv(pk(A)), ik(A), {A, pk(A)}
                  ↪ inv(pk(s)), kdf;

11            B : A, B, s, g, pk(s), pk(A), pk(B), inv(pk(B)), ik(B), kdf;

13            s : A, B, s, g, pk, inv(pk(s)), kdf, ik;

15 where A!=B

17 Actions:

19        B -> s: {B, exp(g, ik(B)), exp(g, OTP), {exp(g, PREKEY)}inv(pk(B)) }
                  ↪ pk(s) # B registers at server

21        s -> A: {B, pk(B)}inv(pk(s)) # s announces B is ready
```

```
23          A -> s : A, B, NA # later in time, A asks to communicate with B

25          s -> A : { A, B, exp(g, ik(B)), exp(g, OTP), {exp(g, PREKEY)}inv(pk(B
               ↪ )), NA }inv(pk(s))

27          A -> B : {A, pk(A)}inv(pk(s)),

29                  {exp(g, OTP), exp(g,EKA),exp(g, ik(A)), exp(g,PREKEY)}inv(pk(
                        ↪ A)),

31                  {| A, B, MSG1 |}kdf(
32                          exp(exp(g,ik(A)), PREKEY), #DH1
33                          exp(exp(g,EKA), ik(B)), #DH2
34                          exp(exp(g,EKA), PREKEY), #DH3
35                          exp(exp(g,OTP), EKA) #DH4
36                          )

38          B -> A : {| B, A, MSG2 |}kdf(
39                          exp(exp(g,ik(A)), PREKEY), #DH1
40                          exp(exp(g,EKA), ik(B)), #DH2
41                          exp(exp(g,EKA), PREKEY), #DH3
42                          exp(exp(g,OTP), EKA) #DH4
43                          )


46  Goals:
47          B authenticates A on exp(g, EKA), exp(g, ik(A))
48          A authenticates s on exp(g, OTP), exp(g, PREKEY), exp(g, ik(B))
49          B authenticates s on exp(g, OTP), exp(g, PREKEY), exp(g, ik(B))
50          MSG1 secret between A,B
51          MSG2 secret between A,B
```

Listing 5.1: X3DH OFMC Model

5.1.2.1 Types section

Here, the parameters of the protocol are defined, e.g. variables, constants, roles, etc. Line 4 defines the participants of the protocol of type **Agent**.

- A and B: Variables indicating Alice and Bob. A variable may be dishonest in an OFMC protocol run.
- s: Constant indicating the server. Constants act as trusted third parties.

Line 5 defines the functions of the protocol.

- kdf: Key derivation function.

- pk: A function to model asymmetric key pairs. A public key for Alice is modeled as pk(A) and the corresponding private key is computed using the internal function inv() as follows inv(pk(A)).
- ik: models the private component of the identity key of a party for the DH calculations.

Alice and Bob have long-term signing keys modeled by the function 'pk()' and long-term identity keys modeled by the function 'ik()'.

Line 6 defines the numbers used in the protocol. Lower-case numbers are constants, while upper-case numbers are random and freshly generated during a protocol run. Some private components of keys are modeled as numbers as they are required to be freshly generated during a protocol run which can not be done using functions.

- g: Public exponent base for DH calculations.
- NA: Alice's Nonce.
- EKA: The private component of Alice's Ephemeral key.
- OTP: The One-time prekey's private component.
- PREKEY: the prekey's private component.
- MSG1 and MSG2: Random numbers used as placeholders for random and fresh messages.

5.1.2.2 Knowledge section

This section of the model defines what knowledge is initially known to each party before the protocol starts. Line 9 defines Alice's knowledge which contains, in addition to the previously defined parameters, Alice's certificate modeled as '{A, pk(A)}inv(pk(s))'. This translates to Alice's identity '**A**' and Alice's public key '**pk(A)**' are signed using the servers private key '**inv(pk(s))**'.

Line 15 holds a 'where' clause that strictly defines A and B as different parties to avoid the scenario where A or B talk to itself.

5.1.2.3 Actions section

The Actions section states the protocol's message sequence between the participating parties for a successful protocol run. Throughout this section messages are referred to according to their line number, e.g message 19 is the message in line number 19 of the model code.

Message 19 represents the registration phase where Bob publishes his public information to the server. The message contains the following:

- exp(g,(XX)): The public DH key of the corresponding private key XX
- {exp(g,PREKEY)}inv(pk(B)): The prekey of Bob signed by Bob's signing key.

The whole message is encrypted for the server for authenticity.

Message 21 is not an actual part of the protocol, however, it is included for the sake of modeling the protocol with the chosen tool. This is because of the asynchronous communication model of OFMC and the strands concept [31].

In message 23, Alice requests to fetch from the server a key bundle to communicate with Bob. A nonce 'NA' is attached for freshness.

Message 25 is the server's response to Alice's key bundle request. The message contains the requested key bundle required to communicate with Bob, as well as Alice's nonce. The whole message is integrity protected by the server's signature.

Message 27 represents the initial message from Alice to Bob. It is composed of the following:

- Alice's certificate for authenticity.
- (Line 29) The Bob's keys which Alice will use to compute SK. All these keys are signed by Alice for integrity protection.
- (Line 31) The initial ciphertext encrypted by the symmetric key SK. The output of the 4 DH operations is input into the KDF resulting in SK.

Message 38 is the last message of the protocol run. It is where Bob shows he is able to compute the same SK.

5.1.2.4 Goals section

This section of the model lists the security goals which the protocol must achieve at the end of a run. There are two types of goals

- Authenticity: When a party wants to make sure a certain message has been generated and sent by the other party. In this model, Alice and Bob authenticate each others' public keys. Bob authenticates the server on the keys he received in message 27 from Alice, to be sure that the keys are forwarded by Alice from the server without tampering.
- Secrecy: The requirement for a message to be secret only between some parties.

5.1.3 Difference from specification

5.1.3.1 Registration Phase

The specification was not specific about how to securely publish Bob's keys to the server and the security of the connection between the server and the communicating parties. Therefore, the registration message from Bob to the server is encrypted by the server's public key.

5.1.3.2 Use of Signatures

The specification discouraged the use of signatures as they reduce deniability which is a desired feature in the messaging application setting which the protocol is intended for. However, this model used signatures.

a) Signature are used in responses from the server to assure integrity of messages to receivers.

b) Alice and Bob have an additional key pair for signing other than the identity key pairs. They are used to sign parts of messages which were vulnerable to attacks by intruders and would break the protocol's security. For example, Bob's keys which Alice used to compute SK in message 27.

5.1.3.3 Use of Certificates

Alice used a certificate signed by the server to authenticate itself in-band as opposed to the specification which ruled the authentication between the parties as a necessary but out of scope operation.

5.2 Double Ratchet Algorithm

The Double Ratchet is a cryptographic asynchronous message exchange algorithm that provides high security to the communicating parties. The algorithm enables the exchange of encrypted messages based on a shared secret key between the two parties where each message is encrypted with its specific ephemeral key. The double ratchet algorithm is a combination of two ratchet constructions which provide enhanced security properties. The first outer ratchet is inherited from Off-the-Record (OTR)'s asymmetric ratchet to benefit from its future secrecy property which is obtained through the use of ephemeral key exchanges. Coupled by an inner symmetric ratchet for forward secrecy, the algorithm was formed and was formerly named Axolotl Ratchet. Additionally, KDF chains are a core concept of the algorithm. This section goes through the methodology of the algorithm, its features, and its security properties.

5.2.1 KDF Chain

As mentioned in section 2.7, a KDF is a function that takes as input a key and an input and produces a cryptographically secure hash output that is random-like. A KDF chain is a series of connected KDFs where one output key of a KDF is a KDF key input of a succeeding KDF. Figure 5.2 shows an illustration for a KDF chain that takes three external inputs and produces three random output keys.
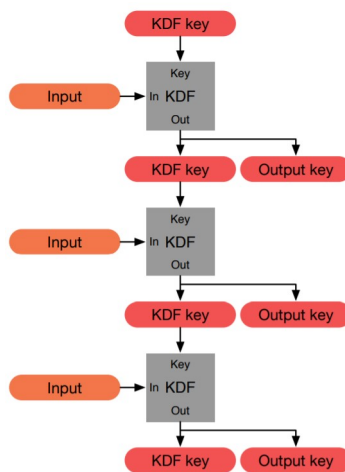
Figure 5.2: A 3 input KDF chain [32].

The algorithm session has three KDF chains: Root chain, sending chain, and receiving chain.

Each chain is advanced whenever its relevant ratchet performs a step. More details about ratchet steps and their relation to the KDF chains are discussed in later sections.

KDF chains have a set of characteristics:

- *Resilience:* An adversary who does not know the KDF keys perceives the output keys as random. Even if the opponent has complete control over the KDF inputs, this is still valid.

- *Forward secrecy:* An adversary who discovers the KDF key at some point in the future cannot distinguish between past output keys and random.

- *Future secrecy:* If future inputs have contributed adequate randomness, future output keys seem random to an adversary who learns the KDF key at some point in the future.

5.2.2 Symmetric-key Ratchet

The sending and receiving chains are constructed of symmetric-key ratchets where Alice's sending chain is equivalent to Bob's receiving chain and vise versa. Essentially, a symmetric-key ratchet is a KDF chain with a constant input through out the chain steps. However, unlike with random input, a constant input does not provide future secrecy.

In a symmetric-key ratchet, the KDF key illustrated in figure 5.2 is the chain key, while a KDF's output key is a unique message key. For a sending chain, a message key is used to encrypt the outgoing message. On the other side, a message key in the receiving chain is used to decrypt the incoming message. The process of calculating the next chain and message keys is an advance in the sending/receiving chain. This is referred to as a *symmetric ratchet step*.

Message keys are not re-introduced into the KDF chain, therefore, they can be discarded without proposing any security risk to earlier or later ratchet outputs. Also, it is possible to store message keys to handle out-of-order messages on the receiving end. Storing message keys introduce a risk to only their respective messages. However, a compromise of a chain key can lead to further compromise of future chain keys as future chain keys rely on previous ones.

5.2.3 Diffie-Hellman Ratchet

The DH ratchet provides the future secrecy property. When paired with the symmetric ratchet, the combination addresses the symmetric ratchet lack of future secrecy. Each party generates a DH key pair known as their *ratchet key pair*. The parties exchange their public keys within message headers with which each can compute an equivalent DH secret output using their own private key that is equivalent to the public key they sent. The process of generating a new DH output when receiving a new public key is referred to as a *DH ratchet step*. Accordingly, a passive adversary that compromises one party's private key at a point in time is incapable of deducing the upcoming DH outputs due to the use of newly generated key pair for each ratchet step.
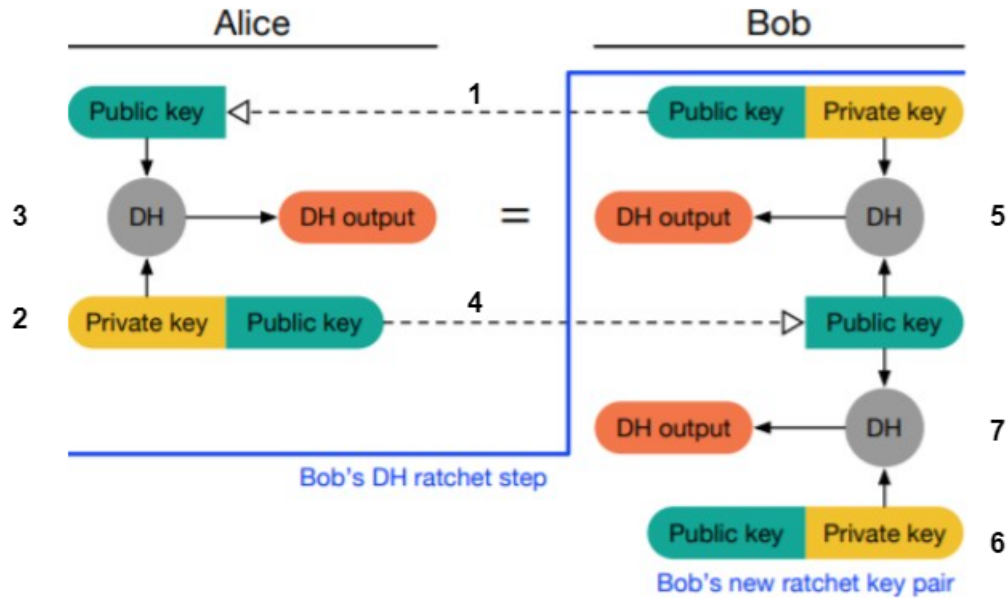
Figure 5.3: Bob DH Ratchet step. Figure reproduced from [32].

Next, we discuss an algorithm run where a DH ratchet is advanced to create two different DH secret outputs between Alice and Bob. Bob is assumed to be the algorithm initiator. Figure 5.3 serves as visual aid for our discussion where we proceed in steps as labeled in the diagram.

- *Step 1:* Bob generates his first ratchet key pair to send the initial message to Alice. The initial message header contains Bob's public key.

- *Step 2:* Upon receiving Bob's public key, Alice generates her ratchet key pair.

- *Step 3:* Alice performs a DH calculation between her private key and Bob's public key generating her side of the DH output.

- *Step 4:* Alice advertises to Bob her public portion of her ratchet key pair that was used to generate the DH output.

- *Step 5:* After obtaining Alice's ratchet public key, Bob performs a DH calculation between it and his current ratchet private key resulting in Bob's copy of the shared DH output.

- *Step 6:* Furthermore, Bob generates a new ratchet key pair to generate the next DH output.

- *Step 7:* Bob performs a DH calculation between hit new ratchet private key and Alice's public key generating a new DH output.

The steps described above are Bob's DH ratchet step as Bob was the initiator of the ratchet step by advertising his public key. Similarly, Alice's ratchet step starts by sending her public key to Bob and concludes after performing the same steps mentioned above but with the roles reversed.

Linking the algorithm to the messaging context, the DH outputs represent the sending and receiving chains' root keys. The party advancing its ratchet, say Bob, results in Alice generating the sending chain root key and Bob generating the receiving chain root key, and vice versa.

However, the description so far is simplified. The algorithm is augmented by a KDF chain to improve resilience and future secrecy. The KDF chain key is a shared secret between both parties while the KDF inputs are the DH outputs from the DH ratchet. Every step through the KDF chain results in a new kdf root key and a sending/receiving chain root key. So a full DH ratchet step consists of updating the root KDF chain twice generating both a sending and a receiving chain key. Figure 5.4 illustrates a full DH ratchet step.
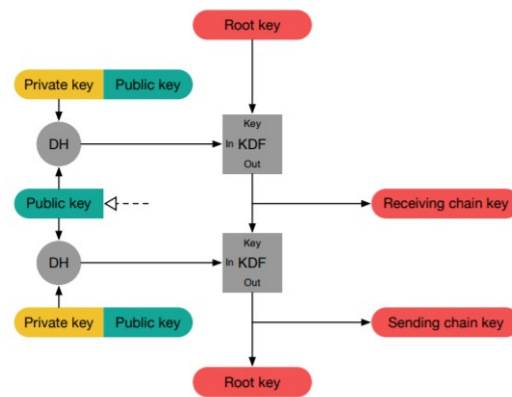
Figure 5.4: A full DH Ratchet step [32].

5.2.4 Double Ratchet

5.2.5 Out-of-order Messages

5.2.6 Header Encryption

5.2.7 Integration with X3DH

5.2.8 Formal Verification

5.2.9 Post-Quantum Security

6 Implementation

7 Evaluation

8 Conclusion and Future Work

Summarize the thesis and provide a outlook on future work.

Bibliography

[1] *2017 Passenger Summary - Annual Traffic Data*. Jan. 2019. URL: https://aci.aero/data-centre/annual-traffic-data/passengers/2017-passenger-summary-annual-traffic-data/.

[2] Boeing. *Commercial Market Outlet 2020 – 2039*. Tech. rep. 2020. URL: https://www.boeing.com/commercial/market/commercial-market-outlook/.

[3] ReadID. *Which countries have ePassports?* 2020. URL: https://www.readid.com/blog/countries-epassports (visited on 08/07/2021).

[4] Lisa Angiolelli-meyer, Jeremy Stokes, and Peter Smallridge. *Automated Border Control Implementation Guide*. Tech. rep. December. IATA, 2015. URL: https://aci.aero/Media/183a7715-da8f-4664-a82c-36f8b961ef28/FiQQnQ/About%20ACI/Priorities/Facilitation/ABC_Implementation_Guide_2nd_Edition.pdf.

[5] Ruggero Donida Labati et al. "Biometric recognition in automated border control: a survey". In: *ACM Computing Surveys (CSUR)* 49.2 (2016), pp. 1–39.

[6] Stefan Marksteiner et al. "An overview of wireless IoT protocol security in the smart home domain". In: *2017 Internet of Things Business Models, Users, and Networks*. IEEE. 2017, pp. 1–8.

[7] Eric Rescorla and Tim Dierks. *The Transport Layer Security (TLS) Protocol Version 1.2*. RFC 5246. Aug. 2008. DOI: 10.17487/RFC5246. URL: https://rfc-editor.org/rfc/rfc5246.txt.

[8] Wentao Shang et al. "Challenges in IoT networking via TCP/IP architecture". In: *NDN Project* (2016).

[9] Eric Rescorla and Nagendra Modadugu. *Datagram Transport Layer Security Version 1.2*. RFC 6347. Jan. 2012. DOI: 10.17487/RFC6347. URL: https://rfc-editor.org/rfc/rfc6347.txt.

[10] Zach Shelby, Klaus Hartke, and Carsten Bormann. *The Constrained Application Protocol (CoAP)*. RFC 7252. June 2014. DOI: 10.17487/RFC7252. URL: https://rfc-editor.org/rfc/rfc7252.txt.

[11] Henrik Nielsen et al. *Hypertext Transfer Protocol – HTTP/1.1*. RFC 2616. June 1999. DOI: 10.17487/RFC2616. URL: https://rfc-editor.org/rfc/rfc2616.txt.

[12] *Bootstrapping*. URL: https://en.wikipedia.org/wiki/Bootstrapping (visited on 08/08/2021).

[13] "IEEE Standard for Local and metropolitan area networks - Secure Device Identity". In: *IEEE Std 802.1AR-2009* (2009), pp. 1–77. DOI: 10.1109/IEEESTD.2009.5367679.

[14] Mohit Sethi, Behcet Sarikaya, and Dan Garcia-Carrillo. *Secure IoT Bootstrapping: A Survey*. Internet-Draft draft-irtf-t2trg-secure-bootstrapping-00. Work in Progress. Internet Engineering Task Force, Apr. 2021. 24 pp. URL: `https://datatracker.ietf.org/doc/html/draft-irtf-t2trg-secure-bootstrapping-00`.

[15] Paul Wouters et al. *Using Raw Public Keys in Transport Layer Security (TLS) and Datagram Transport Layer Security (DTLS)*. RFC 7250. June 2014. DOI: `10.17487/RFC7250`. URL: `https://rfc-editor.org/rfc/rfc7250.txt`.

[16] 3GPP. *Generic Authentication Architecture (GAA); Generic Bootstrapping Architecture (GBA)*. Technical Specification (TS) 33.220. Release 14. 3rd Generation Partnership Project (3GPP), Dec. 2016. URL: `https://portal.3gpp.org/desktopmodules/Specifications/SpecificationDetails.aspx?specificationId=2280`.

[17] Olaf Bergmann, Stefanie Gerdes, and Carsten Bormann. "Simple keys for simple smart objects". In: *Workshop on Smart Object Security*. Vol. 258. Citeseer. 2012.

[18] Max Pritikin et al. *Bootstrapping Remote Secure Key Infrastructure (BRSKI)*. RFC 8995. May 2021. DOI: `10.17487/RFC8995`. URL: `https://rfc-editor.org/rfc/rfc8995.txt`.

[19] Kent Watsen, Mikael Abrahamsson, and Ian Farrer. *Secure Zero Touch Provisioning (SZTP)*. RFC 8572. Apr. 2019. DOI: `10.17487/RFC8572`. URL: `https://rfc-editor.org/rfc/rfc8572.txt`.

[20] Max Pritikin, Peter E. Yee, and Dan Harkins. *Enrollment over Secure Transport*. RFC 7030. Oct. 2013. DOI: `10.17487/RFC7030`. URL: `https://rfc-editor.org/rfc/rfc7030.txt`.

[21] Carsten Bormann, Brian E. Carpenter, and Bing Liu. *GeneRic Autonomic Signaling Protocol (GRASP)*. RFC 8990. May 2021. DOI: `10.17487/RFC8990`. URL: `https://rfc-editor.org/rfc/rfc8990.txt`.

[22] Jeffrey O Kephart and David M Chess. "The vision of autonomic computing". In: *Computer* 36.1 (2003), pp. 41–50.

[23] Stuart Cheshire and Marc Krochmal. *DNS-Based Service Discovery*. RFC 6763. Feb. 2013. DOI: `10.17487/RFC6763`. URL: `https://rfc-editor.org/rfc/rfc6763.txt`.

[24] Stuart Cheshire and Marc Krochmal. *Multicast DNS*. RFC 6762. Feb. 2013. DOI: `10.17487/RFC6762`. URL: `https://rfc-editor.org/rfc/rfc6762.txt`.

[25] Ralph Droms. *Dynamic Host Configuration Protocol*. RFC 2131. Mar. 1997. DOI: `10.17487/RFC2131`. URL: `https://rfc-editor.org/rfc/rfc2131.txt`.

[26] Chris Newman and Graham Klyne. *Date and Time on the Internet: Timestamps*. RFC 3339. July 2002. DOI: `10.17487/RFC3339`. URL: `https://rfc-editor.org/rfc/rfc3339.txt`.

[27] Paul E. Hoffman and Jakob Schlyter. *The DNS-Based Authentication of Named Entities (DANE) Transport Layer Security (TLS) Protocol: TLSA*. RFC 6698. Aug. 2012. DOI: `10.17487/RFC6698`. URL: `https://rfc-editor.org/rfc/rfc6698.txt`.

[28] Alfred J Menezes, Paul C Van Oorschot, and Scott A Vanstone. *Handbook of applied cryptography*. CRC press, 2018.

[29] Moxie Marlinspike and Trevor Perrin. "The x3dh key agreement protocol". In: *Open Whisper Systems* (2016).

[30] Sebastian Mödersheim and Luca Vigano. "The open-source fixed-point model checker for symbolic analysis of security protocols". In: *Foundations of Security Analysis and Design V*. Springer, 2009, pp. 166–194.

[31] Sebastian Mödersheim. *Protocol security verification tutorial*. 2018.

[32] Trevor Perrin and Moxie Marlinspike. *The Double Ratchet Algorithm*. Nov. 2016. URL: `https://www.signal.org/docs/specifications/doubleratchet/doubleratchet.pdf`.

[33] Kent Watsen, Russ Housley, and Sean Turner. *Conveying a Certificate Signing Request (CSR) in a Secure Zero Touch Provisioning (SZTP) Bootstrapping Request*. Internet-Draft draft-ietf-netconf-sztp-csr-12. Work in Progress. Internet Engineering Task Force, Dec. 2021. 36 pp. URL: `https://datatracker.ietf.org/doc/html/draft-ietf-netconf-sztp-csr-12`.

A BRSKI vs. SZTP

A.1 Protocol Comparison

	BRSKI	SZTP
Standardization	RFC 8995	RFC 8572
Related RFCs	Voucher artifact [RFC 8366]	
Aim	Results in the pledge storing a root certificate sufficient for verifying the registrar identity. The installed trust anchor can be used for later certificate enrollment protocols (typically, EST)	Without any manual interference beyond physical placement, securely update the boot image, commit an initial configuration, and execute arbitrary scripts to address auxiliary needs.
Communication channels covered by protocol	Pledge ↔ Registrar ↔ MASA	Device ↔ Owner (↮ MASA: not protocol inherent)
Remote/Local bootstrapping sources support	remote (/.wellknown/...) and local	
Device bootstrap sources	Domain Registrar	Removable storage, DNS server, DHCP server, or Bootstrap server
Protocol initiator	Pledge (Device)	
Functionality support (M): Mandatory (O): Optional	• (M) Pledge-Registrar Discovery • (M) Pledge: polling • (M) if EST following BRSKI: CSR attributes retrieval request • (M) MASA: voucher issuance • (M) MASA: voucher renewal • (M) MASA voucher audit log • (O) Manufacturer: Ownership tracking	• (M) Device: polling • (M) if Bootstrap server is used: provide redirect information and/or onboarding information • (M) DHCP/DNS server: can provide redirect information only due to technical limitations • (M) MASA: voucher issuance

	BRSKI	**SZTP**
Device initial state	• IDevID • Manufacturer installed trust anchor(s) associated with the manufacturer's MASA	• IDevID • Optional: TLS client cert & related intermediate certs • Trust anchors to validate ownership voucher (signed by manufacturer) • List of well-known bootstrap servers • Trust anchors to authenticate configured well-known bootstrap servers
Discovery of bootstrap sources	Yes, mDNS/ GRASP	Only through redirections from device supported bootstrap sources
Device authentication	IDevID	IDevID
Device authorization	• A specific device (serial number) from a specific vendor • A specific device type or a specific vendor	based on device's serial number
Bootstrap source authentication	Intially, Provisional TLS	Initially, Provisional TLS if no TA available
Enrollment protocol integration	(R) EST	Conveying CSR in SZTP draft [33]
Bootstrap data	• Voucher • LDevID	• Redirect information (auxilary) • Onboarding information: boot image, configuration, post-config scripts, voucher, owner certificate

	BRSKI	SZTP
Bootstrapping data protection	• Voucher signed by manufacturer • LDevID signed by domain CA • No additional encryption (only TLS encryption in Transit)	• Trusted channel: may be signed and/or encrypted • Untrusted channel: signed and may be encrypted
Owner voucher-request time	• Nonced: in-band • Nonceless: Out-of-band	• nonced: in-band • Nonceless: owner-manufacturer enrollment phase
Acceptance of device by Domain	Checking voucher and its presence in the MASA audit-log	Checking the voucher
Determining MASA to contact	URI in IDevID or manual confiruration of registrar	Out of scope
Progress reports	Yes, voucher status telemetry	Yes, only to trusted servers
Timeliness	• Nonceless vouchers: expiry time • Nonced vouchers: nonces (and expiry time)	
Revocation checks	• No revocation for vouchers • certificate revocation checks only, depending on pledge capabilities	

	BRSKI	SZTP
Ownership transfer	Yes, by voucher issuance	• Out of scope • (Owner ¡-¿ MASA communication is not inherent to the protocol, but ownership transfer is possible through new vouchers by the MASA)
Updatable Manufacturer Trust Anchors (before bootstrap process)	out-of-scope	Out of scope (through a verifiable process, such as a software upgrade using signed software images)
Transport protocol	HTTP (or CoAP) / TLS1.2+	HTTP/TLS
Required Cryptographic algorithms	None	None
Domain specific configuration provisioning to device	out of scope	yes

A.2 Terminology Comparison

BRSKI		SZTP	
Pledge	The unconfigured device to be bootstrapped	Device	The unconfigured device to be bootstrapped
Manufacturer	The entity that created the device	Manufacturer	• The entity that created the device • Through out the RFC, the term generally covers the services provided by the manufacturer without explicitly referring to them, like, MASA.
Domain	The set of entities, belonging to the owner of the device, that share a common local trust anchor. This includes the proxy, registrar, Domain Certificate Authority, etc..	Owner	• The person or organization that purchased or otherwise owns a device. • The term is used through out the RFC to represent the the sub-entities the owner might be using within their domain, like registrar, proxy, or CA, without explicitly referring to them.
Registrar	A representative of the domain that is configured, to decide whether a new device is allowed to join the domain.		

	BRSKI		SZTP
Proxy	• A domain entity that helps the pledge join the domain. A join proxy facilitates communication for devices that find themselves in an environment where they are not provided connectivity until after they are validated as members of the domain. • The pledge is unaware that they are communicating with a proxy rather than directly with a registrar.		
Voucher	RFC 8366	Ownership voucher	RFC8366